

I will be investigating the impact of SARS-CoV-2, or covid-19, on consumer uncertainty.

- cci - Consumer Confidence Index (CCI) - a monthly survey by The Conference Board that measures how optimistic or pessimistic consumers are regarding financial situations, both current and expected. Administered via survey. Leading indicator reflecting U.S. economic conditions, major purchases, consumer view of economy, business conditions, and labor market currently and over the next six months. Two parts of this index is the current expectations and future expectations
-1985=100
-- cci > 100: consumers are optimistic = more spending
-- cci < 100: consumers are optimistic = less spending
-- cci < 100: consumers are optimistic = less spending

Released on the last Tuesday of every month.

- cci_cur - CCI Current Expectations component
- cci_exp - CCI Future Expectations component
- umsent - Michigan Consumer Sentiment Index (MCSI) -- a monthly report of consumer confidence levels in the U.S. conducted by the UMICH. Survey results collected via telephone interviews. Also a lead indicator. Differs from CCI in that CCI places more weight on employment and labor market while MCSI focuses primarily on households and future expectations is at 12 months instead of 6 months.
- same index number scale

Released on the second Friday of each month.

- umsent_cur - MCSI Current Expectations component
- umsent_exp - MCSI Future Expectations component
- ism_man - ISM Manufacturing Index or Manufacturing Purchasing Managers' Index (PMI) -- reflects the demand level of goods by the amount of ordering activity from factories. Index of new orders, production, employment, supplier deliveries, and inventories.
-- Manufacturing PMI > 50 = expanding manufacturing segment compared to last month
-- Manufacturing PMI < 50 = contracting " "

This is released on the first business day of each month, so can influence CCI and MCSI.

- ism_non - ISM Non-Manufacturing Index or Non-Manufacturing Purchings Managers' Index (PMI) -
- surveys purchasing and supply executives, capturing 15 different (service) industries. Indexes business activity, new orders, employment trends, inventories, and prices.
-- Finance and Insurance, Agriculture, Retail Trade, Utilities, Educational Services, etc.

Released on the third business day of the month.

OLS ASSUMPTIONS:

- [1] The regression model is linear in the coefficients and the error term.
- [2] The error term has a population mean of zero.
- [3] All independent variables are uncorrelated with the error term.
- [4] Observations of the error term are uncorrelated with each other.
- [5] The error term has a constant variance (no heteroscedasticity).
- [6] No independent variable is a perfect linear function of other explanatory variables.
- [7] The error term is normally distributed.

```
In [2]: # libraries
import numpy as np
import statsmodels.api as sm
import pandas as pd
import statsmodels.formula.api as smf
from statsmodels.tsa.ar_model import AutoReg
from statsmodels.tsa.stattools import adfuller
from statsmodels.compat import lzip
import statsmodels.stats.api as sms
from sklearn.linear_model import LinearRegression
from statsmodels.stats.diagnostic import het_white
from statsmodels.graphics.api import interaction_plot, abline_plot
from statsmodels.stats.anova import anova_lm
from statsmodels.graphics.tsaplots import plot_pacf
from scipy import stats
import statsmodels
import matplotlib.pyplot as plt
from matplotlib import pyplot as plt
import seaborn as sns
```

```
In [5]: # data
df = pd.read_csv('cur-us_owid_covid_monthly.csv')
pd.set_option('display.max_columns', None)
df = df.dropna()
# df.shape[0]
df.head(5)
```

```
Out[5]:
```

	date	total_cases	new_cases	new_cases_smoothed	total_deaths	new_deaths	positive_rate	strir
366	7/1/97	0	0	0.0	0	0	0.0	
367	8/1/97	0	0	0.0	0	0	0.0	
368	9/1/97	0	0	0.0	0	0	0.0	
369	10/1/97	0	0	0.0	0	0	0.0	
370	11/1/97	0	0	0.0	0	0	0.0	

```
In [4]: df.shape[0]
```

```
Out[4]: 300
```

```
In [124... x_columns = [ 'total_cases', 'cci_lag1', 'umsent' ]
y = df['cci']
```

```
In [125... x = df[x_columns]
x = sm.add_constant(x)
```

```
summary = sm.OLS(y,x).fit()
print(summary.summary())
```

OLS Regression Results

Dep. Variable:	cci	R-squared:	0.963			
Model:	OLS	Adj. R-squared:	0.963			
Method:	Least Squares	F-statistic:	2580.			
Date:	Tue, 08 Nov 2022	Prob (F-statistic):	8.59e-212			
Time:	14:29:20	Log-Likelihood:	-925.69			
No. Observations:	300	AIC:	1859.			
Df Residuals:	296	BIC:	1874.			
Df Model:	3					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-26.0172	2.758	-9.433	0.000	-31.445	-20.589
total_cases	2.569e-07	3.33e-08	7.707	0.000	1.91e-07	3.22e-07
cci_lag1	0.7284	0.024	30.365	0.000	0.681	0.776
umsent	0.5963	0.052	11.435	0.000	0.494	0.699
=====						
Omnibus:	2.881	Durbin-Watson:	1.863			
Prob(Omnibus):	0.237	Jarque-Bera (JB):	3.144			
Skew:	-0.025	Prob(JB):	0.208			
Kurtosis:	3.499	Cond. No.	1.17e+08			
=====						

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.17e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [126... # Breusch-Pagan test
names = ['Lagrange multiplier statistic', 'p-value',
         'f-value', 'f p-value']

# test result
test_result = sms.het_breuschpagan(summary.resid, summary.model.exog)

lzip(names, test_result)
```

```
Out[126]: [('Lagrange multiplier statistic', 8.25091083213001),
          ('p-value', 0.041101004939359746),
          ('f-value', 2.7903767277975646),
          ('f p-value', 0.04077043426737537)]
```

```
In [127... x_columns1 = [ 'total_cases', 'cci_lag1', 'ism_man']
y1 = df['cci']

x1 = df[x_columns1]
x1 = sm.add_constant(x1)
summary1 = sm.OLS(y1,x1).fit()
print(summary1.summary())
```

OLS Regression Results

Dep. Variable:	cci	R-squared:	0.950			
Model:	OLS	Adj. R-squared:	0.949			
Method:	Least Squares	F-statistic:	1859.			
Date:	Tue, 08 Nov 2022	Prob (F-statistic):	1.19e-191			
Time:	14:29:21	Log-Likelihood:	-972.70			
No. Observations:	300	AIC:	1953.			
Df Residuals:	296	BIC:	1968.			
Df Model:	3					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-12.5427	4.051	-3.096	0.002	-20.515	-4.570
total_cases	-3.12e-08	2.94e-08	-1.062	0.289	-8.9e-08	2.66e-08
cci_lag1	0.9623	0.013	72.542	0.000	0.936	0.988
ism_man	0.3058	0.077	3.992	0.000	0.155	0.457
=====						
Omnibus:	18.157	Durbin-Watson:	1.989			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	41.736			
Skew:	-0.247	Prob(JB):	8.65e-10			
Kurtosis:	4.759	Cond. No.	1.47e+08			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.47e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [128... # Breusch-Pagan test1
names = ['Lagrange multiplier statistic', 'p-value',
         'f-value', 'f p-value']

# test result
test_result1 = sms.het_breuschpagan(summary1.resid, summary1.model.exog)

lzip(names, test_result1)
```

```
Out[128]: [('Lagrange multiplier statistic', 16.47238972372982),
          ('p-value', 0.0009071604648665212),
          ('f-value', 5.732336912412169),
          ('f p-value', 0.0007977702237826029)]
```

```
In [129... x_columns2 = [ 'total_cases', 'cci_lag1', 'ism_non']
y2 = df['cci']

x2 = df[x_columns2]
x2 = sm.add_constant(x2)
summary2 = sm.OLS(y2,x2).fit()
print(summary2.summary())
```

OLS Regression Results

Dep. Variable:	cci	R-squared:	0.951			
Model:	OLS	Adj. R-squared:	0.951			
Method:	Least Squares	F-statistic:	1924.			
Date:	Tue, 08 Nov 2022	Prob (F-statistic):	1.01e-193			
Time:	14:29:22	Log-Likelihood:	-967.87			
No. Observations:	300	AIC:	1944.			
Df Residuals:	296	BIC:	1959.			
Df Model:	3					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-22.3721	5.080	-4.404	0.000	-32.370	-12.374
total_cases	-4.261e-08	2.91e-08	-1.464	0.144	-9.99e-08	1.47e-08
cci_lag1	0.9383	0.014	65.254	0.000	0.910	0.967
ism_non	0.5173	0.101	5.114	0.000	0.318	0.716
=====						
Omnibus:	10.619	Durbin-Watson:	1.999			
Prob(Omnibus):	0.005	Jarque-Bera (JB):	19.513			
Skew:	-0.133	Prob(JB):	5.79e-05			
Kurtosis:	4.221	Cond. No.	1.87e+08			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.87e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [130]: # Breusch-Pagan test2
names = ['Lagrange multiplier statistic', 'p-value',
         'f-value', 'f p-value']

# test result
test_result2 = sms.het_breuschpagan(summary2.resid, summary2.model.exog)

lzip(names, test_result2)
```

```
Out[130]: [('Lagrange multiplier statistic', 21.982382621323282),
          ('p-value', 6.578402699373963e-05),
          ('f-value', 7.801406396786069),
          ('f p-value', 4.980828366660121e-05)]
```

```
In [131]: # Checking for Stationarity, making sure that the mean, variance, and autocorrelation st
# doesn't change over time

df_stationarityTest = adfuller(df['cci'], autolag='AIC')
print('p-value: ', df_stationarityTest[1])

p-value: 0.27891037351945935

for vfi -- multicollinearity ref
```

```
In [132]: # VIF - measuring strength of correlation with between the predictors, checking for mult
# between predictors
x_columns = df[['total_cases', 'cci_lag1', 'umsent']]
x = sm.add_constant(x_columns)
y = df['cci']
```

```
In [133]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
df4 = pd.DataFrame()

df4['VIF'] = [variance_inflation_factor(x_columns3.values, i)
              for i in range(x_columns3.shape[1])]
df4['feature'] = x_columns3.columns

df4
```

Out[133]:

	VIF	feature
0	1.071352	total_cases
1	13.727316	cci_lag1
2	13.762078	ism_man

In [134...]

```
df2 = df

x_columns3 = df[['total_cases', 'cci_lag1', 'ism_man']]
x1 = sm.add_constant(x_columns3)
y1 = df['cci']
```

In [135...]

```
from statsmodels.stats.outliers_influence import variance_inflation_factor

df3 = pd.DataFrame()

df3['VIF'] = [variance_inflation_factor(x_columns3.values, i)
              for i in range(x_columns3.shape[1])]
df3['feature'] = x_columns3.columns

df3
```

Out[135]:

	VIF	feature
0	1.071352	total_cases
1	13.727316	cci_lag1
2	13.762078	ism_man

In [136...]

```
df4 = df
x_columns4 = df[['total_cases', 'cci_lag1', 'ism_man']]
x4 = sm.add_constant(x_columns4)
y4 = df['cci']
```

In [137...]

```
from statsmodels.stats.outliers_influence import variance_inflation_factor

df4 = pd.DataFrame()

df4['VIF'] = [variance_inflation_factor(x_columns4.values, i)
              for i in range(x_columns4.shape[1])]
df4['feature'] = x_columns4.columns

df4
```

Out[137]:

	VIF	feature
0	1.071352	total_cases
1	13.727316	cci_lag1
2	13.762078	ism_man

observed value; resid plot

In [152... *# checking the assumption of normality --homoskedasticity*

data skew and model fit

```
from scipy import stats
```

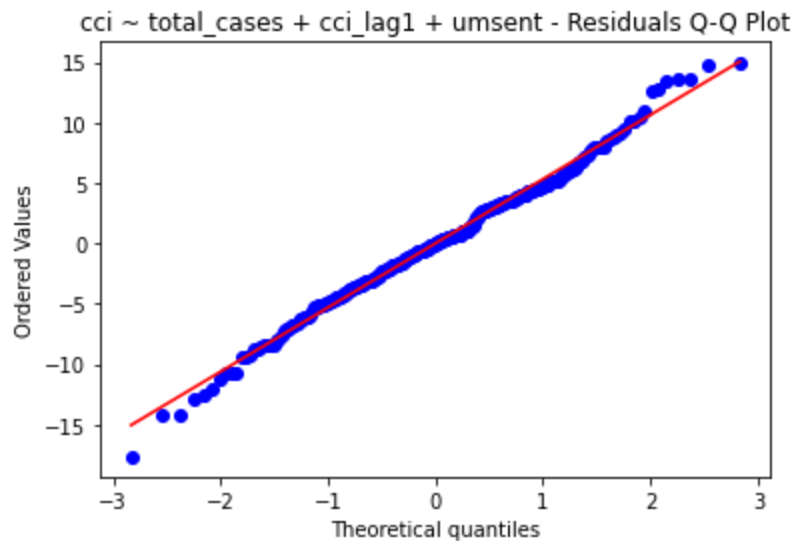
```
stats.probplot(summary.resid, dist="norm", plot= plt)
```

```
plt.title("cci ~ total_cases + cci_lag1 + umsent - Residuals Q-Q Plot")
```

#Saving plot as a png

```
#plt.savefig("Modell_Resid_qqplot.png")
```

Out[152]: Text(0.5, 1.0, 'cci ~ total_cases + cci_lag1 + umsent - Residuals Q-Q Plot')



In [144... **from** scipy **import** stats


```
stats.probplot(summary1.resid, dist="norm", plot= plt)
```

```
plt.title("cci ~ total_cases + cci_lag1 + ism-man - Residuals Q-Q Plot")
```

#Saving plot as a png

```
#plt.savefig("Modell_Resid_qqplot.png")
```

Out[144]: Text(0.5, 1.0, 'cci ~ total_cases + cci_lag1 + ism-man - Residuals Q-Q Plot')




```
In [145... from scipy import stats

stats.probplot(summary2.resid, dist="norm", plot= plt)
plt.title("cci ~ total_cases + cci_lag1 + ism_non - Residuals Q-Q Plot")

#Saving plot as a png
#plt.savefig("Modell_Resid_qqplot.png")
```

Out[145]: Text(0.5, 1.0, 'cci ~ total_cases + cci_lag1 + ism_non - Residuals Q-Q Plot')



```
In [153... table = sm.stats.anova_lm(mo, typ=2)
print(table)
```

	sum_sq	df	F	PR(>F)
total_cases	1687.599639	1.0	59.396999	1.963161e-13
cci_lag1	26196.838522	1.0	922.027691	6.307114e-93
umsent	3715.134343	1.0	130.758402	2.533786e-25
Residual	8410.012276	296.0	NaN	NaN

```
In [147... sy = smf.ols('cci ~ total_cases + cci_lag1 + umsent', df)
mo = sy.fit()
print(mo.summary2())
```


Results: Ordinary least squares

```
=====
Model:                OLS                Adj. R-squared:    0.963
Dependent Variable:   cci                AIC:              1859.3818
Date:                2022-11-08 14:33    BIC:              1874.1970
No. Observations:    300                Log-Likelihood:    -925.69
Df Model:             3                  F-statistic:       2580.
Df Residuals:         296                Prob (F-statistic): 8.59e-212
R-squared:            0.963              Scale:           28.412
=====
```

```
-----
              Coef.   Std.Err.    t    P>|t|    [0.025   0.975]
-----+-----
Intercept    -26.0172   2.7582  -9.4327  0.0000  -31.4454  -20.5891
total_cases    0.0000   0.0000   7.7069  0.0000   0.0000   0.0000
cci_lag1      0.7284   0.0240  30.3649  0.0000   0.6812   0.7756
umsent       0.5963   0.0521  11.4350  0.0000   0.4937   0.6989
=====
```

```
Omnibus:            2.881            Durbin-Watson:        1.863
Prob(Omnibus):       0.237            Jarque-Bera (JB):      3.144
Skew:                -0.025           Prob(JB):              0.208
Kurtosis:            3.499            Condition No.:        116978028
=====
```

* The condition number is large (1e+08). This might indicate strong multicollinearity or other numerical problems.

```
In [148... pg = df.corr()
# pg = pg.to_csv('resid_q-q_plot.csv', encoding='utf-8')
pg
```

Out [148]:

	total_cases	new_cases	new_cases_smoothed	total_deaths	new_deaths	positive_
total_cases	1.000000	0.731918	0.675118	0.983190	0.646423	0.633434
new_cases	0.731918	1.000000	0.916267	0.760064	0.857942	0.829708
new_cases_smoothed	0.675118	0.916267	1.000000	0.698959	0.748111	0.856726
total_deaths	0.983190	0.760064	0.698959	1.000000	0.698337	0.679986
new_deaths	0.646423	0.857942	0.748111	0.698337	1.000000	0.780809
positive_rate	0.633434	0.829708	0.856726	0.679986	0.780809	1.000000
stringency_index	0.595817	0.700117	0.629771	0.699750	0.842179	0.811111
cci	0.094413	0.045618	0.045807	0.104197	0.041404	0.029708
cci_cur	0.152775	0.074558	0.070557	0.151590	0.042262	0.036726
cci_exp	-0.035227	-0.018546	-0.010404	-0.006539	0.031403	0.011111
umsent	-0.329655	-0.279644	-0.248900	-0.326883	-0.262999	-0.297080
umsent_cur	-0.387146	-0.328161	-0.292073	-0.390106	-0.316011	-0.342262
umsent_exp	-0.269706	-0.228921	-0.203722	-0.262853	-0.210001	-0.237267
ism_man	0.253471	0.244688	0.229843	0.289170	0.223233	0.151111
ism_non	0.280734	0.237940	0.208248	0.325448	0.200169	0.129708
cci_lag1	0.098416	0.060226	0.059576	0.104871	0.048653	0.060000
cci_lag2	0.100656	0.080145	0.072448	0.104026	0.082889	0.098337
cci_lag3	0.099427	0.084317	0.073124	0.102212	0.102229	0.105556
cci_lag4	0.098822	0.081383	0.070697	0.102017	0.099910	0.110000
cci_lag5	0.099753	0.089600	0.077625	0.103211	0.103438	0.120000
total_cases_lag1	0.995610	0.678265	0.613095	0.977683	0.597031	0.597031
total_cases_lag2	0.987893	0.667050	0.597711	0.969486	0.559952	0.588889
total_deaths_lag1	0.985699	0.742266	0.680348	0.998647	0.668050	0.668050
total_deaths_lag2	0.986382	0.736626	0.673357	0.995180	0.648733	0.655556

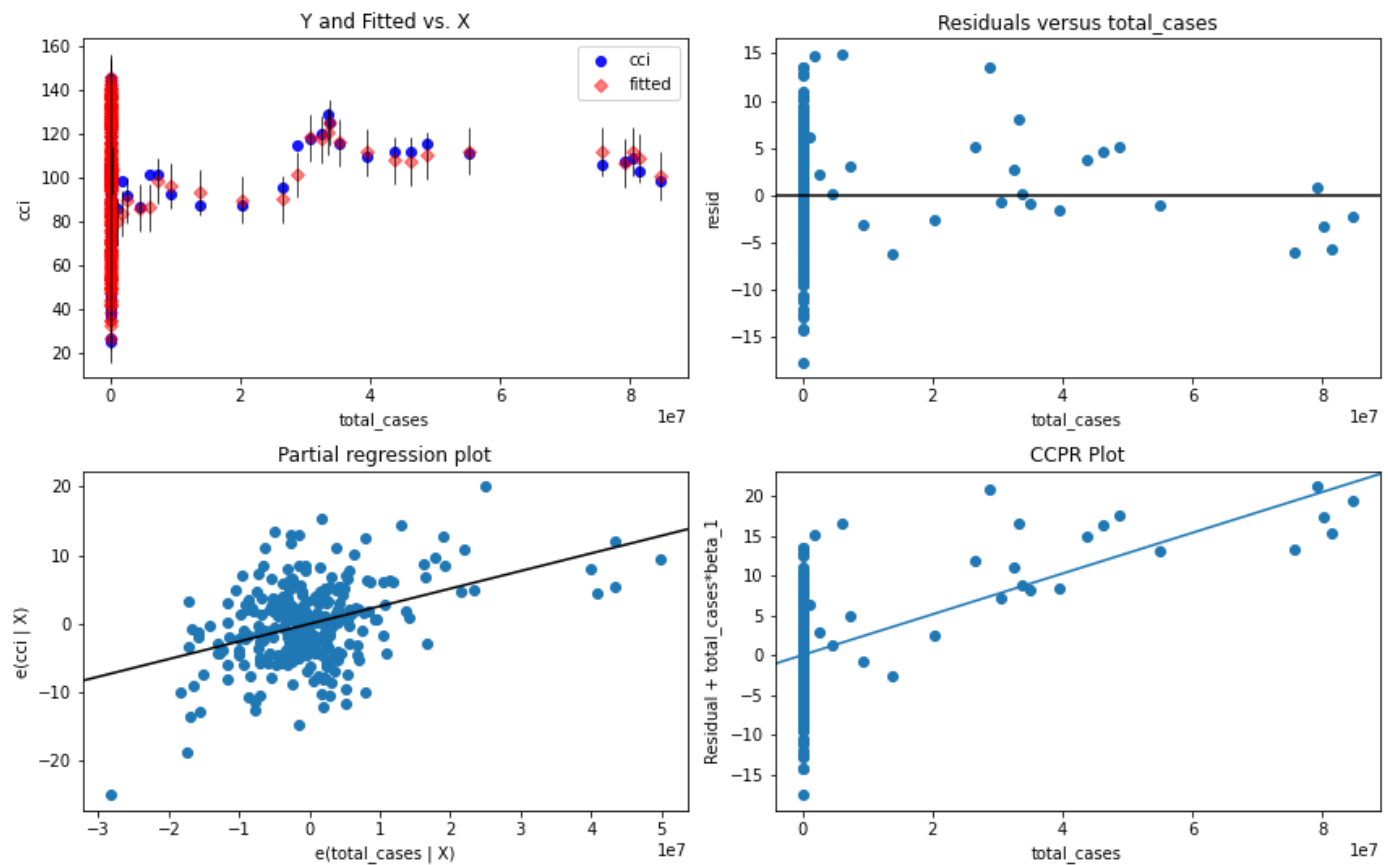
How do I read this?

In [159...]

```
#define figure size
fig =plt.figure(figsize=(12,8))

#produce regression plots
fig = sm.graphics.plot_regress_exog(mo, 'total_cases', fig=fig)

eval_env: 1
```



- 'Y and Fitted vs X' and 'Residuals vs fitted' -- checks linear assumptions of linearity, normality, constant variance (violated-hetero, makes cone shape), and independence.

<https://analyse-it.com/docs/user-guide/fit-model/linear/residual-plot>

- Partial Regression Plot -- checks influential points and linearity
- Partial Regression Plot -- checks influential points, relationship between a regressor and response



by taking into account the other independent variables. The line is if x was highly correlated with any of the other independent variables.

test for heteroskedasticity

```
In [167... # can determine order of AR model
from statsmodels.graphics.tsaplots import plot_pacf
from statsmodels.graphics.tsaplots import plot_acf

pacf = plot_pacf(df['cci'], lags=25)
acf = plot_acf(df['cci'], lags=25)
```


```
/Users/mauricefreese/opt/anaconda3/lib/python3.9/site-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywmm'). You can use this method now by setting method='ywmm'.
warnings.warn(
```




```
In [101... fig = plt.figure(figsize= (10, 10))
ax = fig.add_subplot(111)

normality_plot, stat = stats.probplot(mo.resid, plot= plt, rvalue= True)
ax.set_title("Probability plot of model residual's", fontsize= 20)
ax.set

plt.show()
```



```
In [102... import seaborn as sns
sns.histplot(mo.resid);
```



```
In [103... mu, std = stats.norm.fit(mo.resid)
mu, std
```

```
Out[103]: (5.0401164723249776e-14, 6.290935504613382)
```


```
In [105... fig, ax = plt.subplots()
```

```

# plot the residuals
sns.histplot(x=mo.resid, ax=ax, stat="density", linewidth=0, kde=True)
ax.set(title="Distribution of residuals", xlabel="residual")

# plot corresponding normal curve
xmin, xmax = plt.xlim() # the maximum x values from the histogram above
x = np.linspace(xmin, xmax, 100) # generate some x values
p = stats.norm.pdf(x, mu, std) # calculate the y values for the normal curve
sns.lineplot(x=x, y=p, color="orange", ax=ax)
plt.show()

```



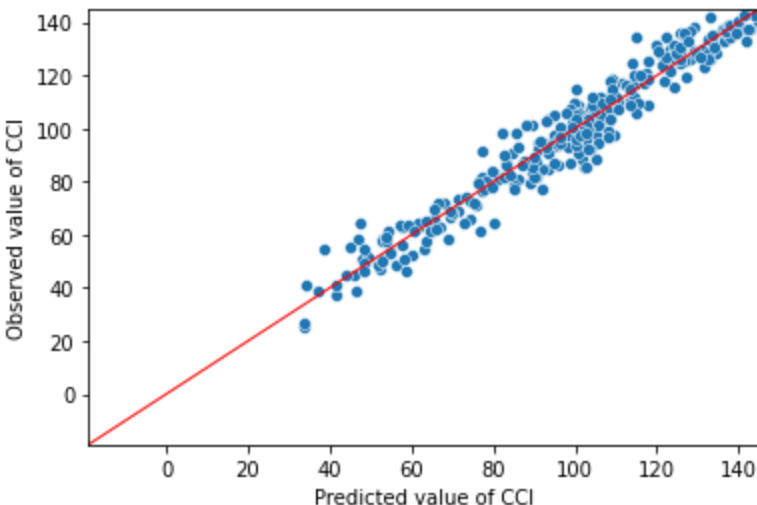
```

In [108... mo.fittedvalues
Y_max = y.max()
Y_min = x.min()

ax = sns.scatterplot(x=mo.fittedvalues, y=y)
ax.set(ylim=(Y_min, Y_max))
ax.set(xlim=(Y_min, Y_max))
ax.set_xlabel("Predicted value of CCI")
ax.set_ylabel("Observed value of CCI")

X_ref = Y_ref = np.linspace(Y_min, Y_max, 100)
plt.plot(X_ref, Y_ref, color='red', linewidth=1)
plt.show()

```



```

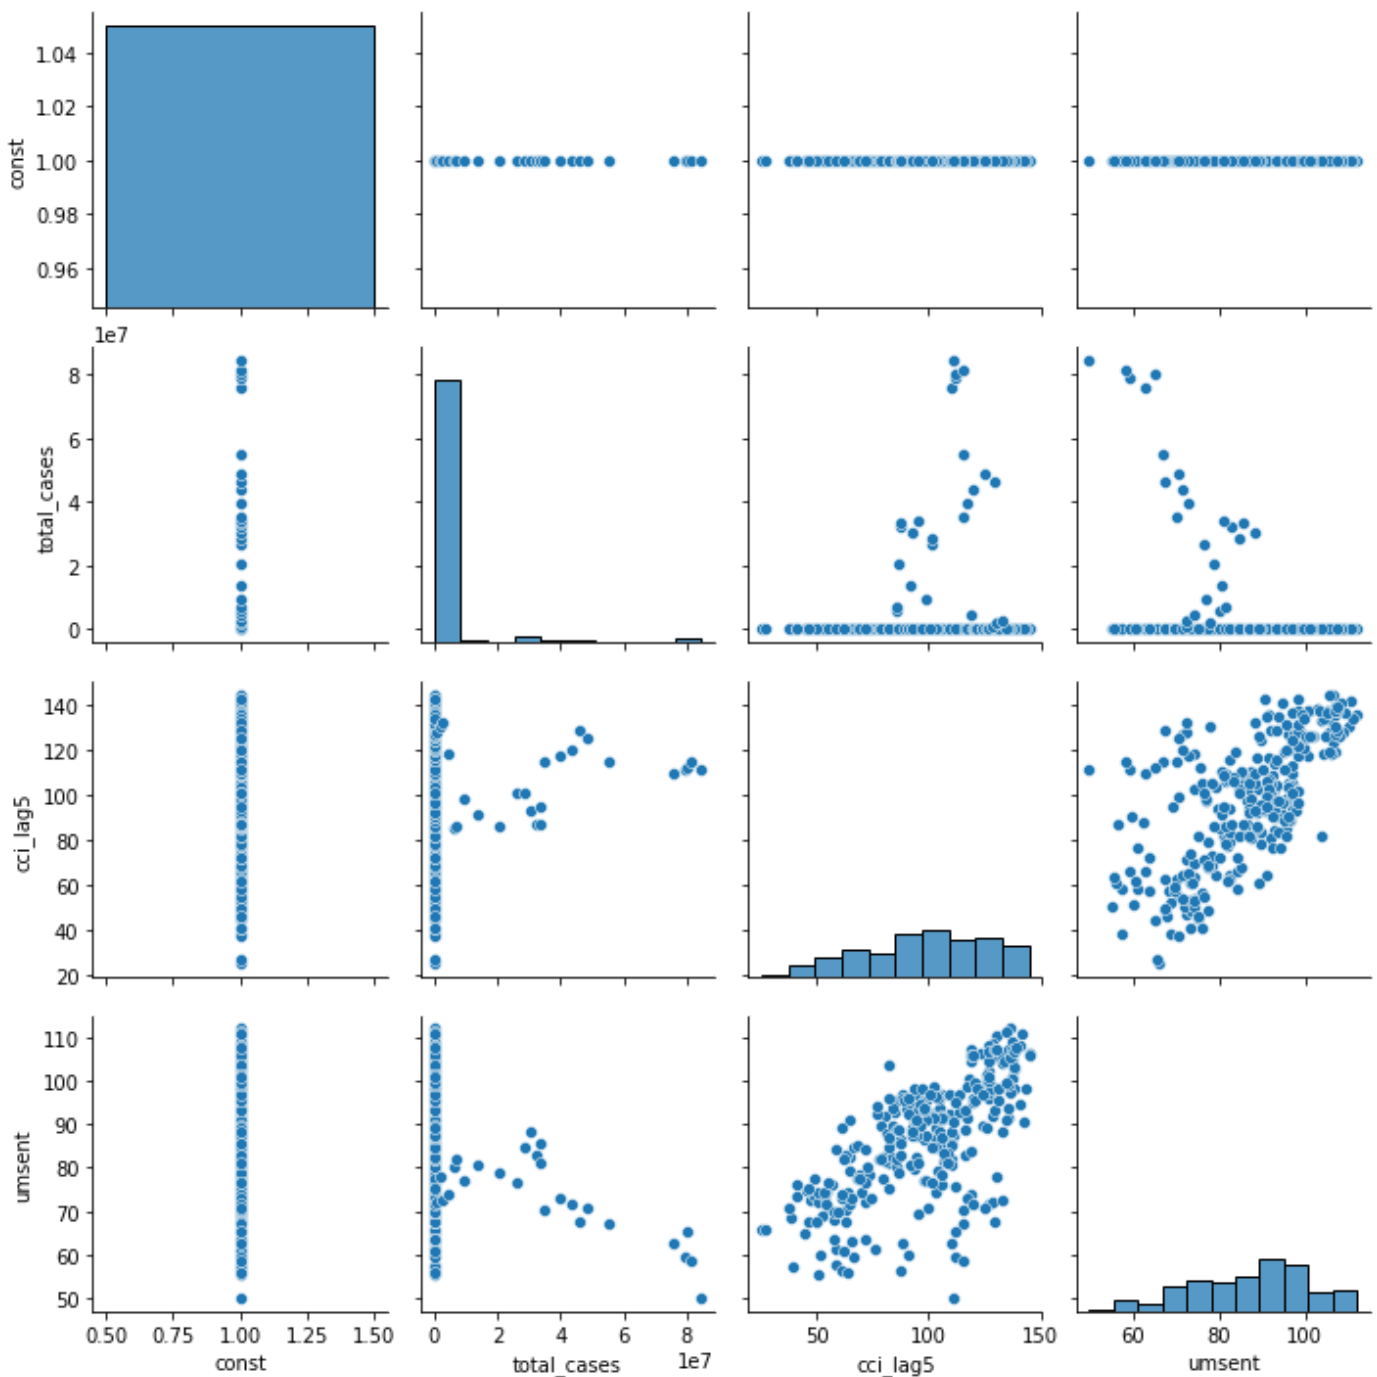
In [117... sns.pairplot(x)

```

```

Out[117]: <seaborn.axisgrid.PairGrid at 0x7fa9c3012100>

```



```
In [122]: # correlation of model
round(x.corr(),3)
```

```
Out[122]:
```

	const	total_cases	cci_lag5	umsent
const	NaN	NaN	NaN	NaN
total_cases	NaN	1.00	0.100	-0.330
cci_lag5	NaN	0.10	1.000	0.677
umsent	NaN	-0.33	0.677	1.000

```
In [128]: from sklearn.cross_decomposition import PLSRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import cross_val_predict

# Define PLS object
pls = PLSRegression(n_components=5)
```

```

# Fit
pls.fit(x_columns, y)

# Cross-validation
y_cv = cross_val_predict(pls, x_columns, y, cv=10)

# Calculate scores
score = r2_score(y, y_cv)
mse = mean_squared_error(y, y_cv)

```

```

-----
ValueError                                Traceback (most recent call last)
Input In [128], in <cell line: 9>()
      6 pls = PLSRegression(n_components=5)
      8 # Fit
----> 9 pls.fit(x_columns, y)
     11 # Cross-validation
     12 y_cv = cross_val_predict(pls, x_columns, y, cv=10)

File ~/opt/anaconda3/lib/python3.9/site-packages/sklearn/cross_decomposition/_pls.py:211
, in _PLS.fit(self, X, Y)
     192 def fit(self, X, Y):
     193     """Fit model to data.
     194
     195     Parameters
     (...)
     208     Fitted model.
     209     """
--> 211     check_consistent_length(X, Y)
     212     X = self._validate_data(
     213         X, dtype=np.float64, copy=self.copy, ensure_min_samples=2
     214     )
     215     Y = check_array(Y, dtype=np.float64, copy=self.copy, ensure_2d=False)

File ~/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py:332, in che
ck_consistent_length(*arrays)
     330 uniques = np.unique(lengths)
     331 if len(uniques) > 1:
--> 332     raise ValueError(
     333         "Found input variables with inconsistent numbers of samples: %r"
     334         % [int(l) for l in lengths]
     335     )

ValueError: Found input variables with inconsistent numbers of samples: [3, 300]

```

In []: